

Data_Cleaning

February 13, 2025

0.1 Exploring the data

Since there is so much data, we need to figure out what the data is and how, if we want to, to combine all the data. We also need to check if anything needs to be cleaned.

```
[22]: import pandas as pd
import numpy as np
import sqlite3

[23]: df_gross = pd.read_csv('../data/bom.movie_gross.csv', index_col=0)
df_budgets = pd.read_csv('../data/tn.movie_budgets.csv', index_col=0)
df_movies = pd.read_csv('../data/tmdb.movies.csv', index_col=0)
df_reviews = pd.read_csv('../data/rt.reviews.tsv', index_col=0, sep='\t',
    ↳encoding='latin-1')
df_info = pd.read_csv('../data/rt.movie_info.tsv', index_col=0, sep='\t')
conn = sqlite3.connect('../data/im.db')
```

This dataset holds the amount of money a movie made domestically and foreignly, the studio associated with the movie and the year it came out.

```
[24]: #df_gross['domestic_gross'] = df_gross['domestic_gross'].map(lambda x: int(x)
    ↳if pd.notnull(x) else x) Not working??
df_gross
```

```
[24]:
```

	studio	domestic_gross	\
title			
Toy Story 3	BV	415000000.0	
Alice in Wonderland (2010)	BV	334200000.0	
Harry Potter and the Deathly Hallows Part 1	WB	296000000.0	
Inception	WB	292600000.0	
Shrek Forever After	P/DW	238700000.0	
...	
The Quake	Magn.	6200.0	
Edward II (2018 re-release)	FM	4800.0	
El Pacto	Sony	2500.0	
The Swan	Synergetic	2400.0	
An Actor Prepares	Grav.	1700.0	
	foreign_gross	year	

title		
Toy Story 3	652000000	2010
Alice in Wonderland (2010)	691300000	2010
Harry Potter and the Deathly Hallows Part 1	664300000	2010
Inception	535700000	2010
Shrek Forever After	513900000	2010
...
The Quake	NaN	2018
Edward II (2018 re-release)	NaN	2018
El Pacto	NaN	2018
The Swan	NaN	2018
An Actor Prepares	NaN	2018

[3387 rows x 4 columns]

This dataset is somewhat similar to the one above, except it has a more specific release date, and also the budget spent on the movie. Upon checking the gross columns, there are no missing boxes, so we can strip the columns and turn them into ints to make them easier to deal with. We will also drop movies before the 2000s as inflation will make the price comparisons not equal.

```
[25]: #df_budgets[['production_budget', 'domestic_gross', 'worldwide_gross']] =
↳df_budgets[['production_budget', 'domestic_gross', 'worldwide_gross']].
↳apply(lambda x: x.replace('$', '').replace(',', ''))
#Remove all symbols and convert to int
df_budgets['production_budget'] = df_budgets['production_budget'].map(lambda x:
↳int(x.replace('$', '').replace(',', '')))
df_budgets['domestic_gross'] = df_budgets['domestic_gross'].map(lambda x: int(x.
↳replace('$', '').replace(',', '')))
df_budgets['worldwide_gross'] = df_budgets['worldwide_gross'].map(lambda x:
↳int(x.replace('$', '').replace(',', '')))
#convert to datetime, remove everything before 2000s
df_budgets['release_date'] = pd.to_datetime(df_budgets['release_date'])
df_budgets = df_budgets[df_budgets['release_date'] > pd.
↳to_datetime('1999-12-31')]
df_budgets
```

```
[25]:   release_date      movie \
id
1    2009-12-18      Avatar
2    2011-05-20  Pirates of the Caribbean: On Stranger Tides
3    2019-06-07      Dark Phoenix
4    2015-05-01  Avengers: Age of Ultron
5    2017-12-15  Star Wars Ep. VIII: The Last Jedi
..    ...
77   2004-12-31    The Mongol King
78   2018-12-31      Red 11
80   2005-07-13  Return to the Land of Wonders
```

```

81 2015-09-29          A Plague So Pleasant
82 2005-08-05          My Date With Drew

```

```

      production_budget  domestic_gross  worldwide_gross
id
1          425000000          760507625          2776345279
2          410600000          241063875          1045663875
3          350000000          42762350           149762350
4          330600000          459005868          1403013963
5          317000000          620181382          1316721747
..          ...          ...          ...
77          7000           900           900
78          7000           0           0
80          5000          1338          1338
81          1400           0           0
82          1100          181041          181041

```

```
[4387 rows x 5 columns]
```

What we can do with this dataset is create a value that measures how much the movie made compared to its production budget to get a simplified value of return.

```

[26]: df_budgets['return_ratio'] = ((df_budgets['worldwide_gross'] -
    ↪df_budgets['production_budget']) / df_budgets['production_budget']).round(2)
df_budgets.sort_values('return_ratio', ascending=False)

```

```

[26]:  release_date      movie  production_budget  domestic_gross  \
id
93  2009-09-25  Paranormal Activity          450000          107918810
80  2015-07-10      The Gallows          100000          22764410
10  2004-05-07    Super Size Me           65000          11529368
82  2005-08-05    My Date With Drew           1100           181041
57  2007-05-16      Once          150000          9445857
..          ...          ...          ...
4   2011-12-31      Tracker          650000           0
81  2016-10-16    Mi America          210000          3330
83  2012-12-31    Infected          210000           0
52  2015-12-11  The Ridiculous 6          6000000           0
14  2015-02-13    Girlhouse          150000           0

```

```

      worldwide_gross  return_ratio
id
93          194183034          430.52
80          41656474          415.56
10          22233808          341.06
82           181041          163.58
57          23323631          154.49
..          ...          ...

```

4	3149	-1.00
81	3330	-1.00
83	0	-1.00
52	0	-1.00
14	0	-1.00

[4387 rows x 6 columns]

This dataset has an interesting column `genre_ids`, which holds arrays of numbers. These numbers presumably can be associated with a dict of some sort that holds what genre it is from the number. It also has a popularity number which isn't obvious what it is, the `vote_average`, which is presumably out of 10, and the vote count. We will drop the `genre_ids` since there doesn't seem to be another table that links it to the actual genres.

```
[27]: df_movies.drop(columns='genre_ids', inplace=True)
df_movies.set_index('id', inplace=True)

df_movies['release_date'] = pd.to_datetime(df_movies['release_date'])
df_movies = df_movies[df_movies['release_date'] > pd.to_datetime('1999-12-31')]
df_movies
```

```
[27]:
```

	original_language	original_title \
id		
12444	en	Harry Potter and the Deathly Hallows: Part 1
10191	en	How to Train Your Dragon
10138	en	Iron Man 2
27205	en	Inception
32657	en	Percy Jackson & the Olympians: The Lightning T...
...
488143	en	Laboratory Conditions
485975	en	_EXHIBIT_84xxx_
381231	en	The Last One
366854	en	Trailer Made
309885	en	The Church

	popularity	release_date \
id		
12444	33.533	2010-11-19
10191	28.734	2010-03-26
10138	28.515	2010-05-07
27205	27.920	2010-07-16
32657	26.691	2010-02-11
...
488143	0.600	2018-10-13
485975	0.600	2018-05-01
381231	0.600	2018-10-01
366854	0.600	2018-06-22
309885	0.600	2018-10-05

	title	vote_average	\
id			
12444	Harry Potter and the Deathly Hallows: Part 1	7.7	
10191	How to Train Your Dragon	7.7	
10138	Iron Man 2	6.8	
27205	Inception	8.3	
32657	Percy Jackson & the Olympians: The Lightning T...	6.1	
...	
488143	Laboratory Conditions	0.0	
485975	_EXHIBIT_84xxx_	0.0	
381231	The Last One	0.0	
366854	Trailer Made	0.0	
309885	The Church	0.0	

	vote_count
id	
12444	10788
10191	7610
10138	12368
27205	22186
32657	4229
...	...
488143	1
485975	1
381231	1
366854	1
309885	1

[26398 rows x 7 columns]

This dataset holds a large amount of reviews with the text, and the score associated with it. This one seems to have many missing rating numbers, some not even being numbers. The fresh section shows this came from rotten tomatoes.

```
[28]: df_reviews
```

```
[28]:
```

	review	rating	fresh	\
id				
3	A distinctly gallows take on contemporary fina...	3/5	fresh	
3	It's an allegory in search of a meaning that n...	NaN	rotten	
3	... life lived in a bubble in financial dealin...	NaN	fresh	
3	Continuing along a line introduced in last yea...	NaN	fresh	
3	... a perverse twist on neorealism...	NaN	fresh	
...	
2000	The real charm of this trifle is the deadpan c...	NaN	fresh	
2000		NaN	1/5	rotten
2000		NaN	2/5	rotten

2000		NaN	2.5/5	rotten
2000		NaN	3/5	fresh

	critic	top_critic		publisher	date
id					
3	PJ Nabarro	0	Patrick Nabarro	November 10,	2018
3	Annalee Newitz	0	io9.com	May 23,	2018
3	Sean Axmaker	0	Stream on Demand	January 4,	2018
3	Daniel Kasman	0	MUBI	November 16,	2017
3	NaN	0	Cinema Scope	October 12,	2017
...
2000	Laura Sinagra	1	Village Voice	September 24,	2002
2000	Michael Szymanski	0	Zap2it.com	September 21,	2005
2000	Emanuel Levy	0	EmanuelLevy.Com	July 17,	2005
2000	Christopher Null	0	Filmcritic.com	September 7,	2003
2000	Nicolas Lacroix	0	Showbizz.net	November 12,	2002

[54432 rows x 7 columns]

This dataset has a lot of info in it. The genre, director, writer, dates, runtime, studio, to name a few which might not be found in the other datasets.

```
[29]: df_info['box_office'] = df_info['box_office'].map(lambda x: int(str(x).
↪replace(',',' ')) if pd.notnull(x) else x)
df_info
```

```
[29]: synopsis rating \
```

id		rating
1	This gritty, fast-paced, and innovative police...	R
3	New York City, not-too-distant-future: Eric Pa...	R
5	Illeana Douglas delivers a superb performance ...	R
6	Michael Douglas runs afoul of a treacherous su...	R
7		NaN NR
...
1996	Forget terrorists or hijackers -- there's a ha...	R
1997	The popular Saturday Night Live sketch was exp...	PG
1998	Based on a novel by Richard Powell, when the l...	G
1999	The Sandlot is a coming-of-age story about a g...	PG
2000	Suspended from the force, Paris cop Hubert is ...	R

	genre	director \
id		
1	Action and Adventure Classics Drama	William Friedkin
3	Drama Science Fiction and Fantasy	David Cronenberg
5	Drama Musical and Performing Arts	Allison Anders
6	Drama Mystery and Suspense	Barry Levinson
7	Drama Romance	Rodney Bennett
...

1996	Action and Adventure Horror Mystery and Suspense	NaN
1997	Comedy Science Fiction and Fantasy	Steve Barron
1998	Classics Comedy Drama Musical and Performing Arts	Gordon Douglas
1999	Comedy Drama Kids and Family Sports and Fitness	David Mickey Evans
2000	Action and Adventure Art House and Internation...	NaN

	writer	theater_date	\
id			
1	Ernest Tidyman	Oct 9, 1971	
3	David Cronenberg Don DeLillo	Aug 17, 2012	
5	Allison Anders	Sep 13, 1996	
6	Paul Attanasio Michael Crichton	Dec 9, 1994	
7	Giles Cooper	NaN	
...	
1996	NaN	Aug 18, 2006	
1997	Terry Turner Tom Davis Dan Aykroyd Bonnie Turner	Jul 23, 1993	
1998	NaN	Jan 1, 1962	
1999	David Mickey Evans Robert Gunter	Apr 1, 1993	
2000	Luc Besson	Sep 27, 2001	

	dvd_date	currency	box_office	runtime	studio
id					
1	Sep 25, 2001	NaN	NaN	104 minutes	NaN
3	Jan 1, 2013	\$	600000.0	108 minutes	Entertainment One
5	Apr 18, 2000	NaN	NaN	116 minutes	NaN
6	Aug 27, 1997	NaN	NaN	128 minutes	NaN
7	NaN	NaN	NaN	200 minutes	NaN
...
1996	Jan 2, 2007	\$	33886034.0	106 minutes	New Line Cinema
1997	Apr 17, 2001	NaN	NaN	88 minutes	Paramount Vantage
1998	May 11, 2004	NaN	NaN	111 minutes	NaN
1999	Jan 29, 2002	NaN	NaN	101 minutes	NaN
2000	Feb 11, 2003	NaN	NaN	94 minutes	Columbia Pictures

[1560 rows x 11 columns]

The rotten tomato datasets may be one we decide to not use. Not only are the columns missing a lot of values, neither datasets have titles that we can tie the data to. Another option would be to instead just attribute the data to studios, but even the studio data is pretty sparse and missing a lot of data.

We need to figure out the structure of this database first. There are quite a few tables and the key that connects them is unknown. There might be data that isn't shown in the other datasets, such as principals, movie_akas, and known_for. So for now, we will explore those since the other tables likely have redundant data.

```
[30]: pd.read_sql(
      """
```

```
SELECT name
FROM sqlite_master
WHERE type='table'
""" , conn)
```

```
[30]:          name
0    movie_basics
1      directors
2     known_for
3    movie_akas
4  movie_ratings
5      persons
6    principals
7      writers
```

It seems like in these databases, people are not referred to by name, instead by an id that will link them to another table. In any case, this table shows the people who worked on a movie, their roles, and potentially the characters name they played.

```
[31]: pd.read_sql(
      """
      SELECT *
      FROM principals
      """ , conn)
```

```
[31]:      movie_id  ordering  person_id  category  job \
0      tt0111414         1    nm0246005    actor    None
1      tt0111414         2    nm0398271  director    None
2      tt0111414         3    nm3739909  producer  producer
3      tt0323808        10    nm0059247    editor    None
4      tt0323808         1    nm3579312  actress    None
...
1028181  tt9692684         1    nm0186469    actor    None
1028182  tt9692684         2    nm4929530     self    None
1028183  tt9692684         3  nm10441594  director    None
1028184  tt9692684         4    nm6009913    writer  writer
1028185  tt9692684         5  nm10441595  producer  producer
```

```
      characters
0      ["The Man"]
1              None
2              None
3              None
4      ["Beth Boothby"]
...
1028181  ["Ebenezer Scrooge"]
1028182  ["Herself", "Regan"]
1028183              None
```



```
1028184          None
1028185          None
```

```
[1028186 rows x 6 columns]
```

This table seems to be for movies that have different titles since they're in a different language.

```
[32]: pd.read_sql(
      """
      SELECT *
      FROM movie_akas
      """, conn)
```

```
[32]:
```

	movie_id	ordering		title	region \
0	tt0369610	10		BG	
1	tt0369610	11		Jurashikku warudo	JP
2	tt0369610	12	Jurassic World: 0	Mundo dos Dinossauros	BR
3	tt0369610	13		0 Mundo dos Dinossauros	BR
4	tt0369610	14		Jurassic World	FR
...
331698	tt9827784	2		Sayonara kuchibiru	None
331699	tt9827784	3		Farewell Song	XWW
331700	tt9880178	1		La atención	None
331701	tt9880178	2		La atención	ES
331702	tt9880178	3		The Attention	XWW

	language	types	attributes	is_original_title
0	bg	None	None	0.0
1	None	imdbDisplay	None	0.0
2	None	imdbDisplay	None	0.0
3	None	None	short title	0.0
4	None	imdbDisplay	None	0.0
...
331698	None	original	None	1.0
331699	en	imdbDisplay	None	0.0
331700	None	original	None	1.0
331701	None	None	None	0.0
331702	en	imdbDisplay	None	0.0

```
[331703 rows x 8 columns]
```

This table relates a person to the movies they worked on.

```
[33]: pd.read_sql(
      """
      SELECT *
      FROM known_for
      """, conn)
```

```
[33]:      person_id  movie_id
0      nm0061671  tt0837562
1      nm0061671  tt2398241
2      nm0061671  tt0844471
3      nm0061671  tt0118553
4      nm0061865  tt0896534
...
1638255 nm9990690  tt9090932
1638256 nm9990690  tt8737130
1638257 nm9991320  tt8734436
1638258 nm9991320  tt9615610
1638259 nm9993380  tt8743182
```

[1638260 rows x 2 columns]

1 Printing out the rest of the data

```
[34]: pd.read_sql(
      """
      SELECT *
      FROM movie_basics
      """, conn)
```

```
[34]:      movie_id      primary_title \
0      tt0063540      Sunghursh
1      tt0066787  One Day Before the Rainy Season
2      tt0069049  The Other Side of the Wind
3      tt0069204      Sabse Bada Sukh
4      tt0100275  The Wandering Soap Opera
...
146139  tt9916538      Kuambil Lagi Hatiku
146140  tt9916622  Rodolpho Teóphilo - O Legado de um Pioneiro
146141  tt9916706      Dankyavar Danka
146142  tt9916730      6 Gunn
146143  tt9916754  Chico Albuquerque - Revelações

      original_title  start_year \
0      Sunghursh      2013
1      Ashad Ka Ek Din      2019
2      The Other Side of the Wind      2018
3      Sabse Bada Sukh      2018
4      La Telenovela Errante      2017
...
146139      Kuambil Lagi Hatiku      2019
146140  Rodolpho Teóphilo - O Legado de um Pioneiro      2015
146141      Dankyavar Danka      2013
```

146142		6 Gunn	2017
146143	Chico Albuquerque - Revelações		2013

	runtime_minutes	genres
0	175.0	Action, Crime, Drama
1	114.0	Biography, Drama
2	122.0	Drama
3	NaN	Comedy, Drama
4	80.0	Comedy, Drama, Fantasy
...
146139	123.0	Drama
146140	NaN	Documentary
146141	NaN	Comedy
146142	116.0	None
146143	NaN	Documentary

[146144 rows x 6 columns]

```
[35]: pd.read_sql(
      """
      SELECT *
      FROM directors
      """, conn)
```

```
[35]:      movie_id  person_id
0      tt0285252  nm0899854
1      tt0462036  nm1940585
2      tt0835418  nm0151540
3      tt0835418  nm0151540
4      tt0878654  nm0089502
...
291169  tt8999974  nm10122357
291170  tt9001390  nm6711477
291171  tt9001494  nm10123242
291172  tt9001494  nm10123248
291173  tt9004986  nm4993825
```

[291174 rows x 2 columns]

```
[36]: pd.read_sql(
      """
      SELECT *
      FROM persons
      """, conn)
```

```
[36]:      person_id  primary_name  birth_year  death_year  \
0      nm0061671  Mary Ellen Bauder         NaN         NaN
```

1	nm0061865	Joseph Bauer	NaN	NaN
2	nm0062070	Bruce Baum	NaN	NaN
3	nm0062195	Axel Baumann	NaN	NaN
4	nm0062798	Pete Baxter	NaN	NaN
...
606643	nm9990381	Susan Grobes	NaN	NaN
606644	nm9990690	Joo Yeon So	NaN	NaN
606645	nm9991320	Madeline Smith	NaN	NaN
606646	nm9991786	Michelle Modigliani	NaN	NaN
606647	nm9993380	Pegasus Envoyé	NaN	NaN

	primary_profession
0	miscellaneous,production_manager,producer
1	composer,music_department,sound_department
2	miscellaneous,actor,writer
3	camera_department,cinematographer,art_department
4	production_designer,art_department,set_decorator
...	...
606643	actress
606644	actress
606645	actress
606646	producer
606647	director,actor,writer

[606648 rows x 5 columns]

```
[37]: pd.read_sql(
      """
      SELECT *
      FROM writers
      """, conn)
```

[37]:	movie_id	person_id
0	tt0285252	nm0899854
1	tt0438973	nm0175726
2	tt0438973	nm1802864
3	tt0462036	nm1940585
4	tt0835418	nm0310087
...
255868	tt8999892	nm10122246
255869	tt8999974	nm10122357
255870	tt9001390	nm6711477
255871	tt9004986	nm4993825
255872	tt9010172	nm8352242

[255873 rows x 2 columns]

```
[38]: pd.read_sql(
      """
      SELECT *
      FROM movie_ratings
      """, conn)
```

```
[38]:
```

	movie_id	averagerating	numvotes
0	tt10356526	8.3	31
1	tt10384606	8.9	559
2	tt1042974	6.4	20
3	tt1043726	4.2	50352
4	tt1060240	6.5	21
...
73851	tt9805820	8.1	25
73852	tt9844256	7.5	24
73853	tt9851050	4.7	14
73854	tt9886934	7.0	5
73855	tt9894098	6.3	128

[73856 rows x 3 columns]

With some basic cleaning done, and a little quick analysis, we can begin doing our own separate analysis and creating business recommendations

```
[39]: df_gross.to_csv('../data/cleaned_movie_gross.csv', encoding='utf-8')
      df_budgets.to_csv('../data/cleaned_budgets.csv', encoding='utf-8')
      df_movies.to_csv('../data/cleaned_movies.csv', encoding='utf-8')
      df_info.to_csv('../data/cleaned_rt_info.csv', encoding='utf-8')
```